

## Q1: Data Frame

### 1. Read the dataset

```
python
Copy code
import pandas as pd

# Load the Iris dataset
url = "https://github.com/neurospin/pystatsml/raw/master/datasets/iris.csv"
df = pd.read_csv(url)
```

### 2. Print column names

```
python
Copy code
print(df.columns.tolist())
```

### 3. Get numerical columns

```
python
Copy code
numerical_columns = df.select_dtypes(include=['float64', 'int64']).columns.tolist()
print(numerical_columns)
```

### 4. Compute mean of numerical columns grouped by species

```
python
Copy code
stats_table = df.groupby("species")[numerical_columns].mean().reset_index()
print(stats_table)
```

---

## Q2: McCulloch-Pitts (M-P) Neuron

Write a Python program to define the weights and bias for the described M-P neuron:

```
python
Copy code
def mcculloch_pitts(x, y, z):
```

```
weights = [1, 1] # Equal weights for x and y
bias = -1      # Threshold value

# Weighted sum
weighted_sum = weights[0] * x + weights[1] * y + bias

# Activation logic
if x == -1 and y == 1:
    return z # z is the output if condition matches
return -1 # Otherwise
```

Example usage:

```
python
Copy code
print(mcculloch_pitts(-1, 1, 1)) # Outputs 1
print(mcculloch_pitts(0, 1, -1)) # Outputs -1
```

---

### Q3: Data Preparation

#### 1. Load the heart dataset

```
python
Copy code
heart_url = "https://www.kaggle.com/zhaoyingzhu/heartcsv/download"
heart_data = pd.read_csv(heart_url)
```

#### 2. Perform operations

##### a) Find Shape of Data

```
python
Copy code
print(heart_data.shape)
```

##### b) Find Missing Values

```
python
Copy code
```

```
print(heart_data.isnull().sum())
```

#### c) Find Data Types

```
python
```

Copy code

```
print(heart_data.dtypes)
```

#### d) Find Zeros

```
python
```

Copy code

```
zeros_per_column = (heart_data == 0).sum()
```

```
print(zeros_per_column)
```

#### e) Find Mean Age

```
python
```

Copy code

```
mean_age = heart_data['Age'].mean()
```

```
print(mean_age)
```

#### f) Extract Specific Columns and Split Data

```
python
```

Copy code

```
selected_columns = ['Age', 'Sex', 'ChestPain', 'RestBP', 'Chol']
```

```
subset = heart_data[selected_columns]
```

```
# Split into training and testing datasets
```

```
train_data = subset.sample(frac=0.75, random_state=42)
```

```
test_data = subset.drop(train_data.index)
```

```
print(train_data.shape, test_data.shape)
```

### Q4: Matrix and Vector Computations

The matrix  $M$  and vector  $v$  are given as:

$M = \begin{bmatrix} 1 & 2 & 3 \\ 2 & 3 & 0 \\ 1 & 3 & 3 \end{bmatrix}$ ,  $v = \begin{bmatrix} 3 \\ 0 \\ 2 \end{bmatrix}$   $M = 121233303$ ,  $v = 302$

We need to compute:

1.  $M \cdot v$
2.  $v^T \cdot M$
3.  $v^T \cdot v$

Using Python and NumPy:

```
python
Copy code
import numpy as np

# Define M and v
M = np.array([[1, 2, 3], [2, 3, 0], [1, 3, 3]])
v = np.array([3, 0, 2])

# Compute M · v
Mv = np.dot(M, v)

# Compute vT · M
vT_M = np.dot(v.T, M)

# Compute vT · v
vT_v = np.dot(v.T, v)

# Display results
print("M · v =", Mv)
print("vT · M =", vT_M)
print("vT · v =", vT_v)
```

**Results:**

- $M \cdot v = \begin{bmatrix} 9 \\ 6 \\ 9 \end{bmatrix}$
- $v^T \cdot M = \begin{bmatrix} 9 & 15 & 15 \end{bmatrix}$
- $v^T \cdot v = 13$

---

### Q5: Regression Technique

To analyze the temperature dataset, follow these steps:

#### 1. Load the dataset:

```
python
Copy code
import pandas as pd

# Load the temperature data
url = "https://www.kaggle.com/venky73/temperatures-of-india/download"
data = pd.read_csv(url)
print(data.head())
```

#### 2. Preprocess the data:

- Handle missing values (if any).
- Normalize or standardize the data if necessary.

```
python
Copy code
# Check for missing values
print(data.isnull().sum())

# Fill missing values (if any)
data.fillna(data.mean(), inplace=True)
```

#### 3. Fit a regression model:

```
python
Copy code
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import train_test_split

# Define features (X) and target (y)
X = data[['Month']] # Example feature
```

```
y = data['Temperature'] # Target variable

# Split the data
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=42)

# Fit the regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict and evaluate
predictions = model.predict(X_test)
print("Model Coefficients:", model.coef_)
print("Model Intercept:", model.intercept_)
```

#### 4. Plot results (optional):

```
python
Copy code
import matplotlib.pyplot as plt

plt.scatter(X_test, y_test, color='blue', label='Actual')
plt.plot(X_test, predictions, color='red', label='Predicted')
plt.xlabel('Month')
plt.ylabel('Temperature')
plt.legend()
plt.show()
```

## Q6. Classification Technique (Decision Tree for Graduate Admissions)

### A. Data Preprocessing

```
python
Copy code
import pandas as pd
from sklearn.model_selection import train_test_split
```

```
from sklearn.preprocessing import LabelEncoder
from sklearn.tree import DecisionTreeClassifier
from sklearn.metrics import accuracy_score, classification_report

# Load Dataset
url = 'https://www.kaggle.com/mohansacharya/graduate-admissions'
data = pd.read_csv(url)

# Data Cleaning & Transformation
data['Admitted'] = data['Admitted'].apply(lambda x: 1 if x == 'yes' else 0) # Adjust target column
X = data.drop('Admitted', axis=1)
y = data['Admitted']
```

## **B. Data Preparation (Train-Test Split)**

```
python
Copy code
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)
```

## **C. Apply Machine Learning Algorithm**

```
python
Copy code
# Train Decision Tree Classifier
clf = DecisionTreeClassifier(random_state=42)
clf.fit(X_train, y_train)

# Make Predictions
y_pred = clf.predict(X_test)
```

## **D. Model Evaluation**

```
python
Copy code
# Evaluate Model
accuracy = accuracy_score(y_test, y_pred)
```

```
print("Accuracy:", accuracy)
print(classification_report(y_test, y_pred))
```

---

## Q7. Clustering Techniques (Mall Customers Dataset)

### Load Dataset and Preprocess

```
python
Copy code
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
from sklearn.preprocessing import StandardScaler

# Load Dataset
url = 'https://www.kaggle.com/shwetabh123/mall-customers'
data = pd.read_csv(url)

# Select Relevant Features
features = data[['Annual Income', 'Spending Score']]
scaler = StandardScaler()
scaled_features = scaler.fit_transform(features)
```

### Apply K-Means Clustering

```
python
Copy code
# Apply K-Means
kmeans = KMeans(n_clusters=3, random_state=42)
data['Cluster'] = kmeans.fit_predict(scaled_features)

# Visualize Clusters
plt.figure(figsize=(10, 6))
for cluster in range(3):
    plt.scatter(features.iloc[data['Cluster'] == cluster, 0],
```



```
features.iloc[data['Cluster'] == cluster, 1], label=f'Cluster {cluster}')

plt.title('K-Means Clustering of Mall Customers')
plt.xlabel('Annual Income')
plt.ylabel('Spending Score')
plt.legend()
plt.show()
```

---

## Q5. Linear Regression for Month-Wise Temperature

### A. Apply Linear Regression

```
python
Copy code
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, mean_absolute_error, r2_score

# Assuming `temperature_data` is loaded
X = temperature_data[['Month']] # Features
y = temperature_data['Temperature'] # Target

# Train-Test Split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)

# Train Model
lr = LinearRegression()
lr.fit(X_train, y_train)

# Predictions
y_pred = lr.predict(X_test)
```

## B. Performance Assessment

```
python  
Copy code  
# Evaluate Model  
  
mse = mean_squared_error(y_test, y_pred)  
mae = mean_absolute_error(y_test, y_pred)  
r2 = r2_score(y_test, y_pred)  
  
print(f"MSE: {mse}, MAE: {mae}, R2: {r2}")
```

## C. Visualization

```
python  
Copy code  
# Plot Regression Line  
  
plt.scatter(X_test, y_test, color='blue', label='Actual')  
plt.plot(X_test, y_pred, color='red', label='Predicted')  
  
plt.title('Month-Wise Temperature Prediction')  
plt.xlabel('Month')  
plt.ylabel('Temperature')  
plt.legend()  
plt.show()
```

## Q8. Association Rule Learning (Market Basket Optimization)

### A. Data Preprocessing

```
python  
Copy code  
import pandas as pd  
  
from mlxtend.frequent_patterns import apriori, association_rules  
  
# Load Dataset  
  
dataset = pd.read_csv('https://www.kaggle.com/hemantkumar05/market-basket-optimization',  
header=None)
```

```
# Convert Dataset into a Transaction List
```

```
transactions = []
```

```
for i in range(0, len(dataset)):
```

```
    transactions.append([str(dataset.values[i, j]) for j in range(0, dataset.shape[1]) if  
str(dataset.values[i, j]) != 'nan'])
```

## **B. Generate List of Transactions**

```
python
```

```
Copy code
```

```
from mlxtend.preprocessing import TransactionEncoder
```

```
# Transform Transaction Data
```

```
te = TransactionEncoder()
```

```
te_array = te.fit(transactions).transform(transactions)
```

```
df = pd.DataFrame(te_array, columns=te.columns_)
```

## **C. Train Apriori Algorithm**

```
python
```

```
Copy code
```

```
# Apply Apriori
```

```
frequent_itemsets = apriori(df, min_support=0.01, use_colnames=True)
```

```
# Generate Association Rules
```

```
rules = association_rules(frequent_itemsets, metric="lift", min_threshold=1)
```

## **D. Visualize the Rules**

```
python
```

```
Copy code
```

```
# Print and Visualize Rules
```

```
print(rules[['antecedents', 'consequents', 'support', 'confidence', 'lift']])
```

---

## Q9. Clustering Techniques (Mall Customers Dataset)

### A. Apply Preprocessing

```
python
Copy code
from sklearn.preprocessing import StandardScaler

# Load Dataset
data = pd.read_csv('https://www.kaggle.com/shwetabh123/mall-customers')

# Select Spending Score for Clustering
X = data[['Spending Score']]
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

### B. Perform Train-Test Split

(Not required for unsupervised learning; we use the entire dataset.)

### C. Apply Machine Learning Algorithms

#### 1. K-Means Clustering

```
python
Copy code
from sklearn.cluster import KMeans

# Apply K-Means
kmeans = KMeans(n_clusters=3, random_state=42)
data['KMeans_Cluster'] = kmeans.fit_predict(X_scaled)
```

#### 2. Hierarchical Clustering

```
python
Copy code
from scipy.cluster.hierarchy import linkage, dendrogram, fcluster
import matplotlib.pyplot as plt
```

```
# Apply Hierarchical Clustering
Z = linkage(X_scaled, method='ward')
dendrogram(Z)
plt.title('Dendrogram for Hierarchical Clustering')
plt.show()

# Assign Clusters
data['Hierarchical_Cluster'] = fcluster(Z, t=3, criterion='maxclust')
```

#### D. Evaluate Model

```
python
Copy code
# Visualize Clusters (K-Means Example)

plt.scatter(X_scaled[data['KMeans_Cluster'] == 0], [0] * len(X_scaled[data['KMeans_Cluster'] == 0]), label='Cluster 0')

plt.scatter(X_scaled[data['KMeans_Cluster'] == 1], [0] * len(X_scaled[data['KMeans_Cluster'] == 1]), label='Cluster 1')

plt.scatter(X_scaled[data['KMeans_Cluster'] == 2], [0] * len(X_scaled[data['KMeans_Cluster'] == 2]), label='Cluster 2')

plt.legend()
plt.show()
```

#### E. Apply Cross-Validation

Cross-validation is typically not applicable for clustering. Instead, clustering performance can be evaluated using metrics like the silhouette score:

```
python
Copy code
from sklearn.metrics import silhouette_score

silhouette_avg = silhouette_score(X_scaled, data['KMeans_Cluster'])
print("Silhouette Score for K-Means:", silhouette_avg)
```

## Q9. Clustering Techniques (Mall Customers Dataset)

### A. Preprocessing

```
python
Copy code
from sklearn.preprocessing import StandardScaler

# Load the dataset
data = pd.read_csv('https://www.kaggle.com/shwetabh123/mall-customers')

# Select Spending Score for Clustering
X = data[['Spending Score']]

# Standardize the data
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

### B. Apply Machine Learning Algorithms

#### 1. K-Means Clustering

```
python
Copy code
from sklearn.cluster import KMeans

# Apply K-Means
kmeans = KMeans(n_clusters=3, random_state=42)
data['KMeans_Cluster'] = kmeans.fit_predict(X_scaled)
```

#### 2. DBSCAN Clustering

```
python
Copy code
from sklearn.cluster import DBSCAN

# Apply DBSCAN
```

```
dbscan = DBSCAN(eps=0.5, min_samples=5)
data['DBSCAN_Cluster'] = dbscan.fit_predict(X_scaled)
```

### C. Evaluate Models

```
python
Copy code
from sklearn.metrics import silhouette_score

# Evaluate K-Means
kmeans_silhouette = silhouette_score(X_scaled, data['KMeans_Cluster'])
print("Silhouette Score for K-Means:", kmeans_silhouette)

# Evaluate DBSCAN
dbscan_silhouette = silhouette_score(X_scaled, data['DBSCAN_Cluster'], metric='euclidean')
print("Silhouette Score for DBSCAN:", dbscan_silhouette)
```

### D. Cross-Validation

Cross-validation for clustering is non-trivial; use silhouette scores as a proxy for evaluating different numbers of clusters in K-Means.

```
python
Copy code
# Example: Cross-validation for K-Means
for n_clusters in range(2, 6):
    kmeans = KMeans(n_clusters=n_clusters, random_state=42)
    clusters = kmeans.fit_predict(X_scaled)
    silhouette_avg = silhouette_score(X_scaled, clusters)
    print(f"Silhouette Score for {n_clusters} clusters: {silhouette_avg}")
```

---

## Q10. PCA on the Iris Dataset

### A. Describe the Dataset

The Iris dataset contains 150 samples with four features: sepal length, sepal width, petal length, and petal width, classified into three species. Standardization is necessary because PCA is sensitive to the scale of the data.

## B. Standardization

```
python
Copy code
from sklearn.datasets import load_iris
from sklearn.preprocessing import StandardScaler

# Load dataset
iris = load_iris()
X = iris.data
y = iris.target

# Standardize the data
scaler = StandardScaler()
X_standardized = scaler.fit_transform(X)
```

## C. Compute PCA

```
python
Copy code
from sklearn.decomposition import PCA

# Apply PCA
pca = PCA(n_components=4)
X_pca = pca.fit_transform(X_standardized)
```

## D. Principal Component Directions and Correlations

```
python
Copy code
# Explained variance
print("Explained Variance Ratio:", pca.explained_variance_ratio_)

# Principal components
print("Principal Components:\n", pca.components_)
```



### E. Plot Samples Projected into First K PCs

python

Copy code

```
import matplotlib.pyplot as plt

# Plot the first two principal components
plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis', edgecolor='k')
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.title('PCA on Iris Dataset')
plt.colorbar()
plt.show()
```

### F. Color Samples by Their Species

python

Copy code

```
# Visualize with species
species = iris.target_names

plt.scatter(X_pca[:, 0], X_pca[:, 1], c=y, cmap='viridis', edgecolor='k')
plt.xlabel('First Principal Component')
plt.ylabel('Second Principal Component')
plt.title('PCA - Iris Species')
plt.colorbar(ticks=[0, 1, 2], label='Species')
plt.clim(-0.5, 2.5)
plt.show()
```